

First kernel

TOPICS

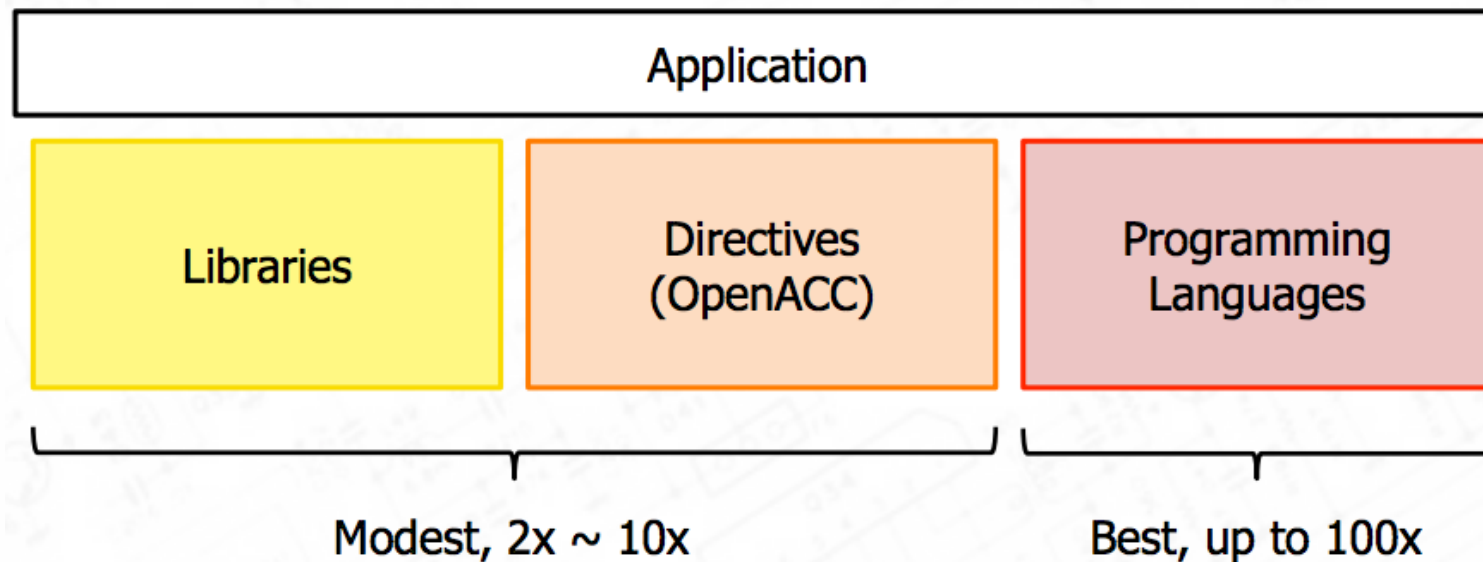
- How to identify the basic terminology used in CUDA (Compute Unified Device Architecture)
- Difference between device (GPU) and host (CPU)
- Code your first application to be executed in the device

Key words: `cudaMalloc`, `cudaFree`, `threaldx`, `cudaMemcpy`, `cudaMemcpyHostToDevice`, `cudaMemcpyDeviceToHost`, `nvcc`.




- GPU is specialized for compute-intensive, highly parallel computation - exactly what graphics rendering is about.
- It is a parallel programming paradigm, the GPU is especially well-suited to address problems that can be expressed as data-parallel computations.
- CUDA, a General-Purpose Parallel Computing Platform and Programming Model
- CUDA comes with a software environment that allows developers to use C as a high-level programming language.

- CUDA runtime system provides an application-programming interface (API) to start working
- API are functions that help users to perform some of the functionalities.
- <http://docs.nvidia.com/cuda/>



■ CUDA ecosystem





Kernel

- CUDA C extends C by allowing the programmer to define C functions, called *kernels*, that, when called, are executed N times in parallel by N different *CUDA threads*, as opposed to only once like regular C functions.
- This kind of function indicates to the compiler that the code will run on the device, and it's callable from the host.

```
__global__ void function(parameter0, parameter1,...,  
parameterN) {  
  
...  
  
}
```

- A call to a `__global__` function is asynchronous, meaning it returns before the device has completed its execution.
- To invoke the kernel function, it is required to use angle brackets `<<<n,n>>>`.

```
function<<<1,1>>>(parameter1, parmeter2,..., parameterN)
```

- The angle brackets indicate how the runtime will launch the device code.



Compilation

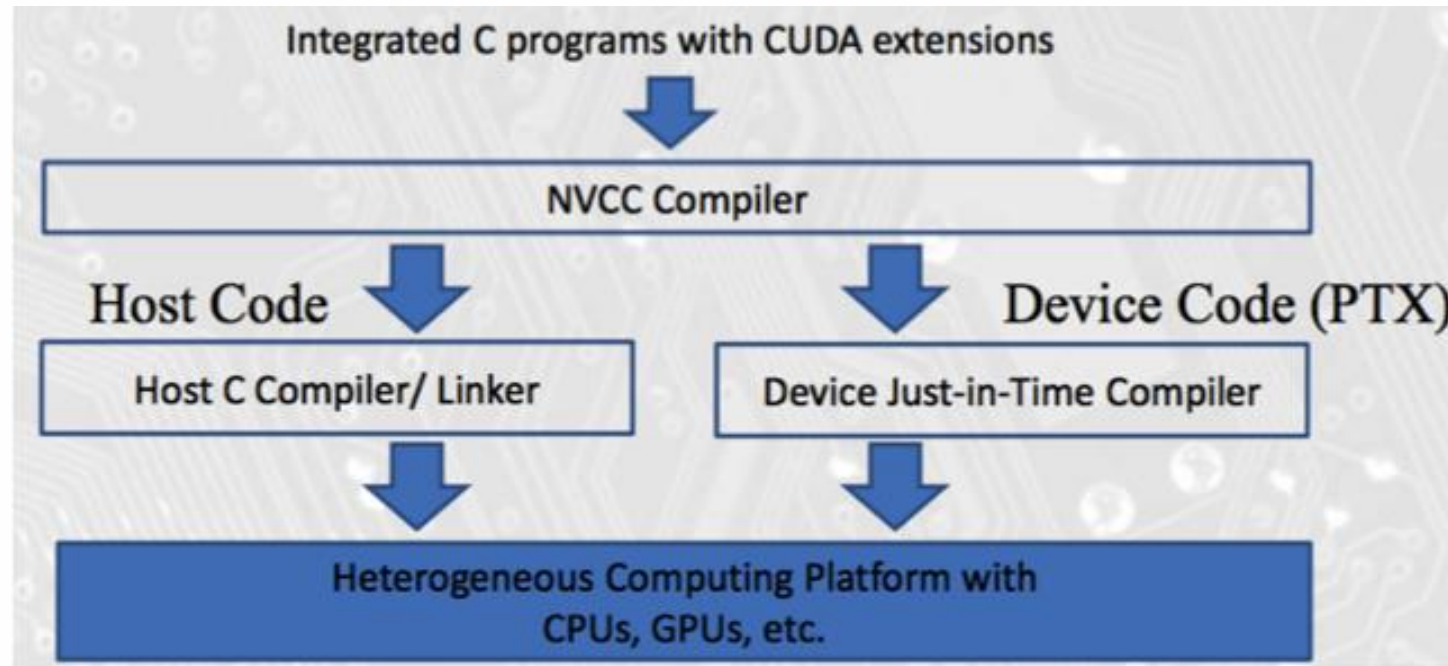
- To compile the example we need to call the Nvidia compiler:

```
$ nvcc source_code.cu -o exe
```

- To run the example:

```
$ ./exe
```

- Kernels can be written using the CUDA instruction set architecture, called PTX.
- **nvcc** is a compiler driver that simplifies the process of compiling C or PTX code: It provides simple and familiar command line options and executes them by invoking the collection of tools that implement the different compilation stages.





Compilation Workflow

- Source files compiled with `nvcc` can include a mix of host code (i.e., code that executes on the host) and device code (i.e., code that executes on the device).
- `nvcc`'s basic workflow consists in separating device code from host code.
 - compiling the device code into an assembly form (*PTX* code) and/or binary form (*cubin* object)
 - and modifying the host code by replacing the `<<<...>>>` syntax by the necessary CUDA C runtime function calls to load and launch each compiled kernel from the *PTX* code and/or *cubin* object.
- The modified host code is output either as C code that is left to be compiled using another tool or as object code directly by letting `nvcc` invoke the host compiler during the last compilation stage.

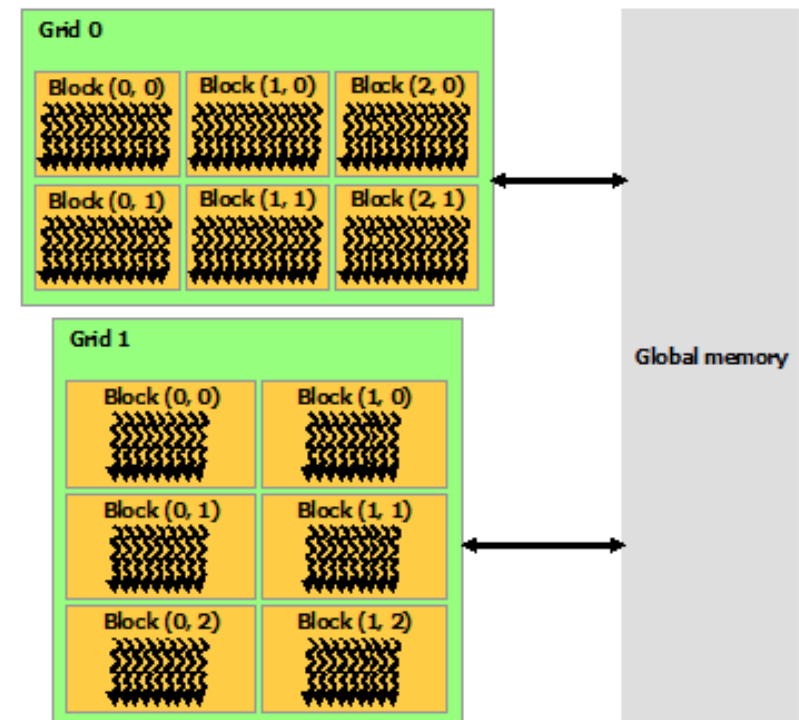



Example

- `01simple_kernel`: This example only creates a kernel to be called, but it doesn't execute any actual instruction in the device.

Memory Hierarchy

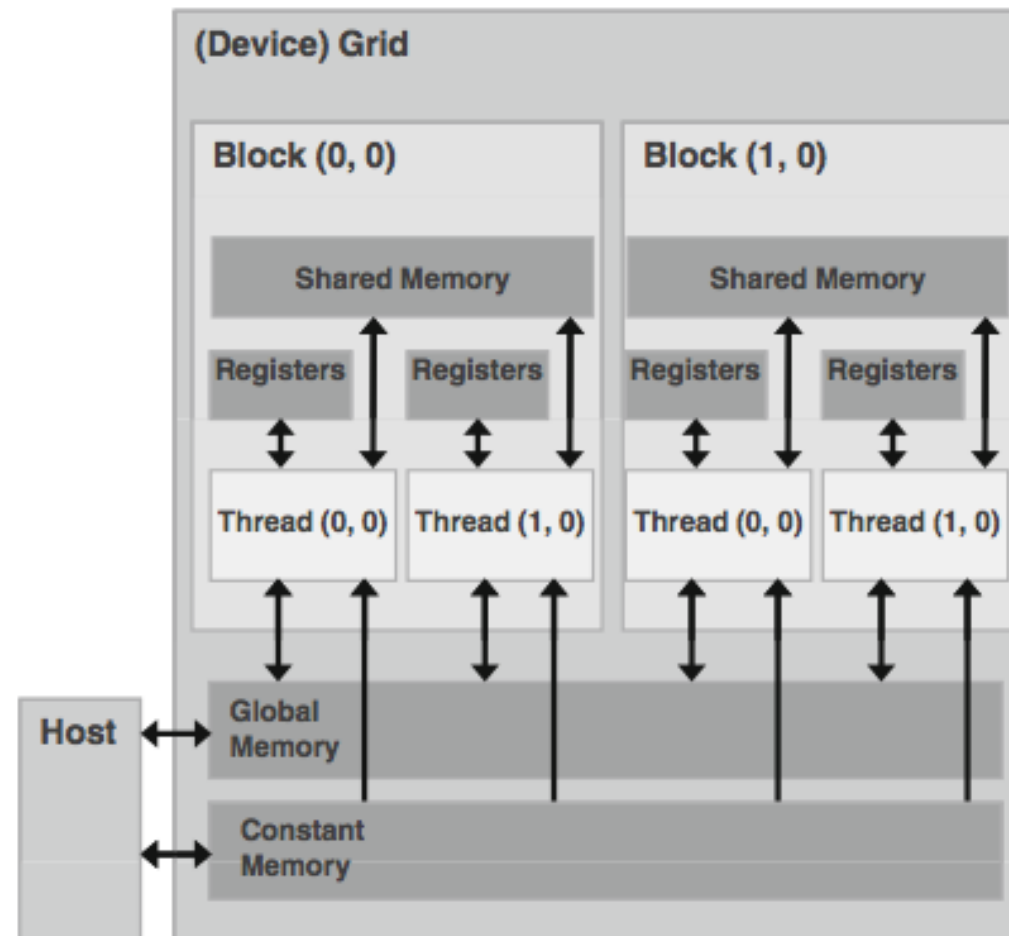
- Each thread has private local memory.
- Each thread block has shared memory visible to all threads of the block and with the same lifetime as the block.
- All threads have access to the same global memory.





CUDA device memory

- Device code can:
 - R/W per-thread registers
 - R/W per-thread local memory
 - R/W per-block shared memory
 - R/W per-grid global memory
 - Read only per-grid constant memory
- Host code can
 - Transfer data to/from per-grid global and constant memories





cudaMalloc

- Besides the kernel function, it is necessary to allocate memory on the device for the parameters that the kernel needs.

```
cudaError_t    cudaMalloc(void** devPtr, size_t size)
```

where:

devPtr - Pointer to allocated device memory

size - Requested allocation size in bytes

- Host pointers can access memory from the host code, and device pointers can access memory from the device code.
- This memory, allocated with `cudaMalloc()` is called linear memory.



cudaFree

- Similar to C programming language, it is necessary to free the memory that is no longer needed, but it must be the one allocated on the device.

```
cudaError_t cudaFree (void * devPtr)
```



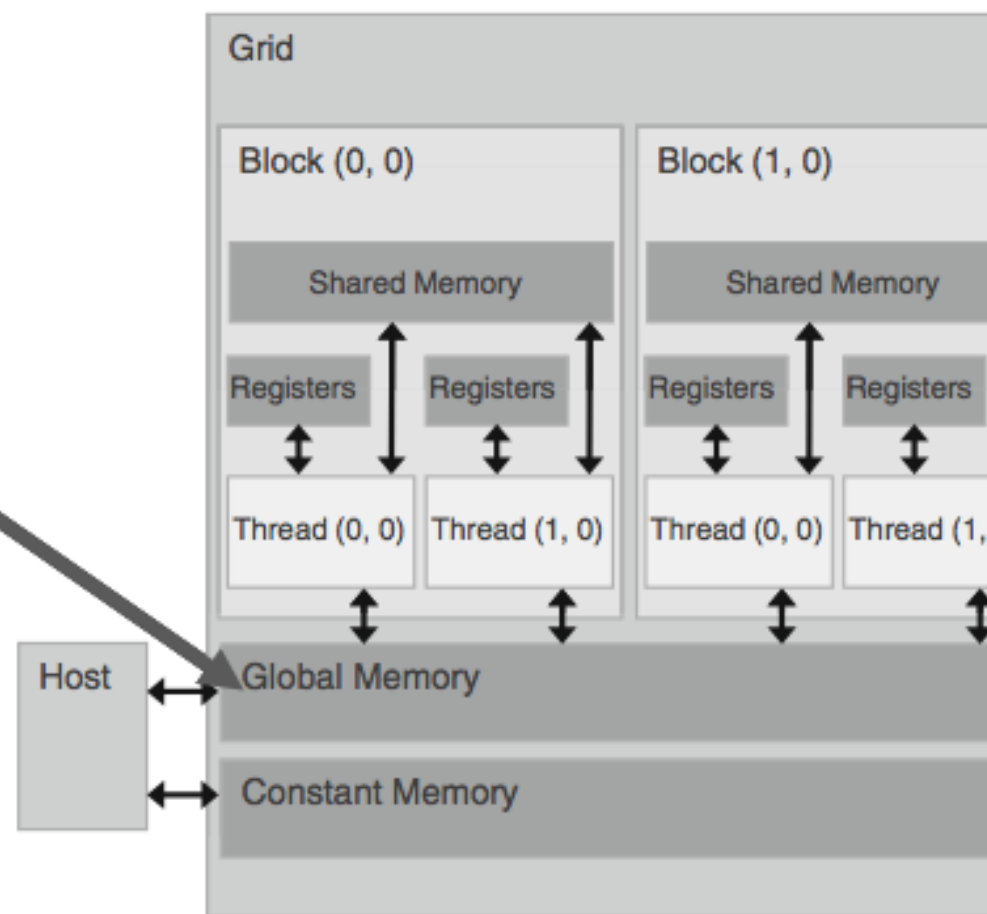
CUDA device memory

`cudaMalloc()`

- Allocates object in the device global memory
- Two parameters
 - **Address of a pointer** to the allocated object
 - **Size of** allocated object in terms of bytes

`cudaFree()`

- Frees object from device global memory
 - Pointer to freed object





cudaMemcpy

- Finally to copy the memory on the device, the *cudaMemcpy()* function is employed.

```
cudaError_t cudaMemcpy( void * dst, const void * src, size_t  
count, enum cudaMemcpyKind kind)
```

where:

dst - Destination memory address

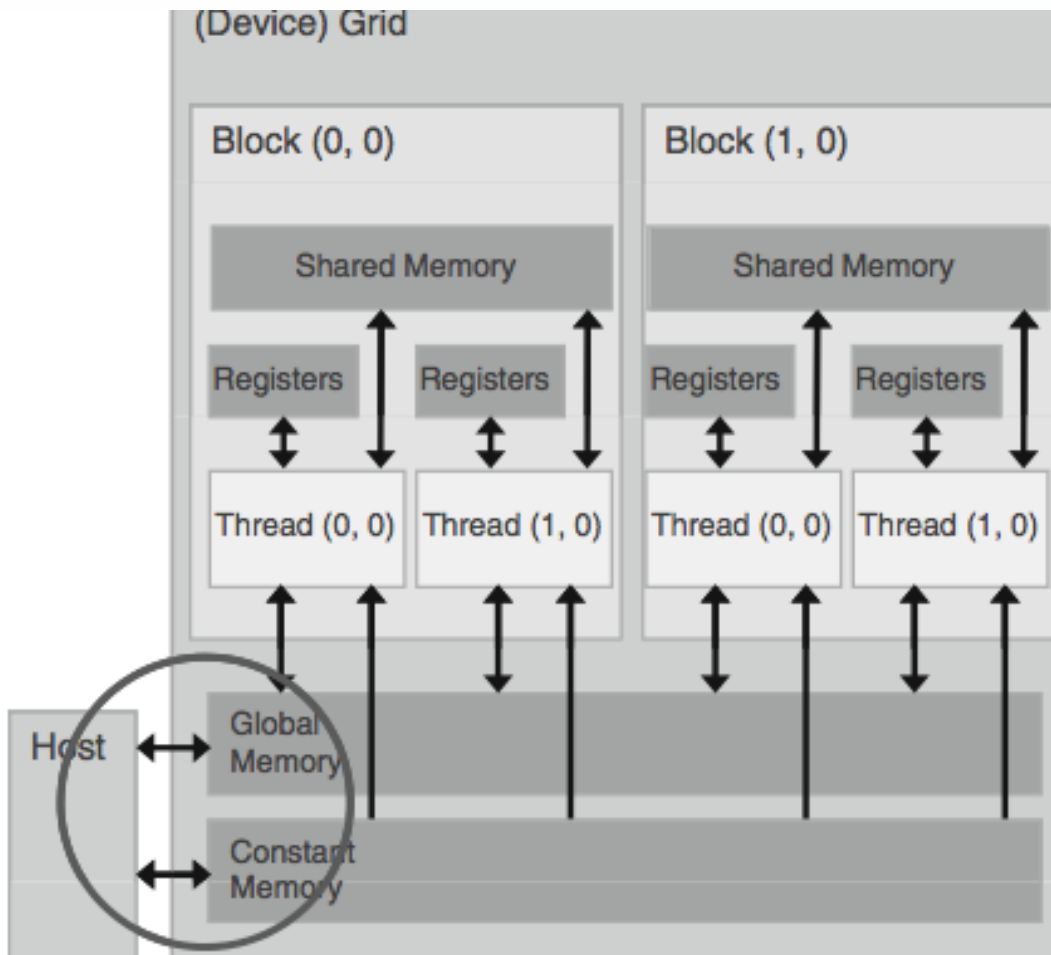
src - Source memory address

count - Size in bytes to copy

kind - Type of transfer (cudaMemcpyHostToDevice,
cudaMemcpyDeviceToHost,)

CUDA device memory

- `cudaMemcpy()`
 - **Memory** data transfer
 - Requires four parameters
 - Pointer to destination
 - Pointer to source
 - Number of bytes copied
 - Type of transfer
 - Host to Host
 - Host to Device
 - Device to Host
 - Device to Device
- Transfer is asynchronous





Example

- `02simple_kernel2`: This code sends three parameters to the kernel and executes the addition operation on them.